

# movielens

September 24, 2021

```
[1]: # import required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: # read the data files and save to variable
data_movie=pd.read_csv('movies.dat', sep='::',
↳names=['MovieID','Title','Genre'], engine='python')
```

```
[3]: data_movie.head()
```

```
[3]:
```

	MovieID	Title	Genre
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

```
[4]: data_movie.shape
```

```
[4]: (3883, 3)
```

```
[5]: data_ratings=pd.read_csv('ratings.dat', sep='::', names=['UserID', 'MovieID',
↳'Rating','Timestamp'], engine='python')
```

```
[6]: data_ratings.head()
```

```
[6]:
```

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

```
[7]: data_ratings.shape
```

```
[7]: (1000209, 4)
```

```
[8]: data_users=pd.read_csv('users.dat', sep='::', names=['UserID', 'Gender', 'Age', 'Occupation', 'Zip-code'], engine='python')
```

```
[9]: data_users.head()
```

```
[9]:  UserID  Gender  Age  Occupation  Zip-code
0      1      F    1         10      48067
1      2      M   56         16      70072
2      3      M   25         15      55117
3      4      M   45          7      02460
4      5      M   25         20      55455
```

```
[10]: data_users.shape
```

```
[10]: (6040, 5)
```

```
[11]: # merge the datasets
Master_Data_tmp=pd.merge(data_users[['UserID', 'Age', 'Gender', 'Occupation']],data_ratings[['UserID', 'MovieID', 'Rating']], on=['UserID'])
```

```
[12]: Master_Data_tmp.head()
```

```
[12]:  UserID  Age  Gender  Occupation  MovieID  Rating
0      1    1      F         10      1193      5
1      1    1      F         10      661      3
2      1    1      F         10      914      3
3      1    1      F         10     3408      4
4      1    1      F         10     2355      5
```

```
[13]: Master_Data_tmp.shape
```

```
[13]: (1000209, 6)
```

```
[14]: # the final merged Master Dataset
Master_Data=pd.merge(Master_Data_tmp,data_movie[['MovieID', 'Title']],on=['MovieID'])
```

```
[15]: Master_Data.head()
```

```
[15]:  UserID  Age  Gender  Occupation  MovieID  Rating  \
0      1    1      F         10      1193      5
1      2   56      M         16      1193      5
2     12   25      M         12      1193      4
3     15   25      M          7      1193      4
4     17   50      M          1      1193      5
```

Title

```

0 One Flew Over the Cuckoo's Nest (1975)
1 One Flew Over the Cuckoo's Nest (1975)
2 One Flew Over the Cuckoo's Nest (1975)
3 One Flew Over the Cuckoo's Nest (1975)
4 One Flew Over the Cuckoo's Nest (1975)

```

```
[16]: Master_Data.shape
```


```
[16]: (1000209, 7)
```

## 0.1 Histogram of User Age distribution

```

[17]: fig = plt.figure(figsize = (10, 5))
counts,_=np.histogram(data_users.Age, bins=[1,18,25,35,45,50,56,100])
plt.pie(counts,labels=["Under 18",
    ↪ "18-24", "25-34", "35-44", "45-49", "50-55", "56+"])
plt.show()

```



## 0.2 User rating of the movie “Toy Story”

```
[18]: ToyStory=Master_Data[Master_Data.Title=='Toy Story (1995)'][['Rating']]
```

```
[19]: ToyStory
```

```
[19]:      Rating
41626      5
41627      4
41628      4
41629      5
41630      5
...      ...
43698      5
43699      5
43700      4
43701      4
43702      3
```

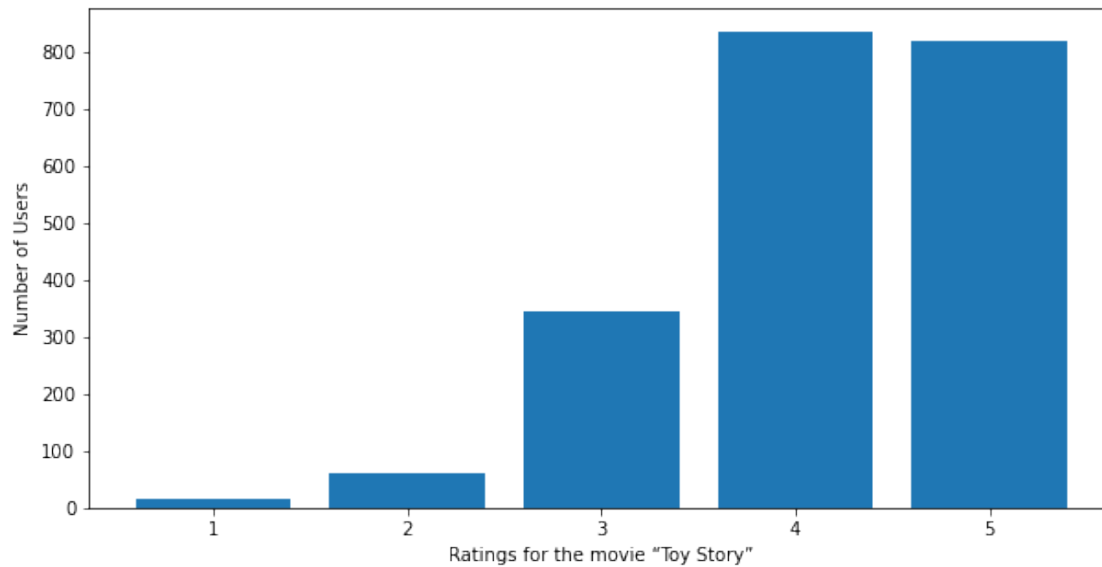
[2077 rows x 1 columns]

```
[20]: rating_count=ToyStory['Rating'].value_counts()
```

```
[21]: print(rating_count)
```

```
4      835
5      820
3      345
2       61
1       16
Name: Rating, dtype: int64
```

```
[22]: # Histogram of User rating of the movie "Toy Story"
fig = plt.figure(figsize = (10, 5))
counts1,_=np.histogram(ToyStory, bins=[1,2,3,4,5,6])
plt.bar([1,2,3,4,5],counts1)
plt.xticks([1,2,3,4,5], [1,2,3,4,5])
plt.xlabel('Ratings for the movie "Toy Story"')
plt.ylabel('Number of Users')
plt.show()
```



### 0.3 Top 25 movies by viewership rating

```
[23]: gr_count=Master_Data.groupby('Title')['Rating'].count().reset_index(name='No. of Ratings')
      gr_mean=Master_Data.groupby('Title')['Rating'].mean().reset_index(name='Avg. Rating')
```

```
[24]: gr_count
```

```
[24]:
```

	Title	No. of Ratings
0	\$1,000,000 Duck (1971)	37
1	'Night Mother (1986)	70
2	'Til There Was You (1997)	52
3	'burbs, The (1989)	303
4	...And Justice for All (1979)	199
...	...	...
3701	Zed & Two Noughts, A (1985)	29
3702	Zero Effect (1998)	301
3703	Zero Kelvin (Kjrlighetens kjtere) (1995)	2
3704	Zeus and Roxanne (1997)	23
3705	eXistenZ (1999)	410

```
[3706 rows x 2 columns]
```

```
[25]: gr_mean
```

```
[25]:
```

	Title	Avg. Ratings
0	\$1,000,000 Duck (1971)	3.027027
1	'Night Mother (1986)	3.371429
2	'Til There Was You (1997)	2.692308
3	'burbs, The (1989)	2.910891
4	...And Justice for All (1979)	3.713568
...	...	...
3701	Zed & Two Noughts, A (1985)	3.413793
3702	Zero Effect (1998)	3.750831
3703	Zero Kelvin (Kjrlighetens kjtere) (1995)	3.500000
3704	Zeus and Roxanne (1997)	2.521739
3705	eXistenZ (1999)	3.256098

[3706 rows x 2 columns]

```
[26]: merged_rating=pd.merge(gr_count,gr_mean,on=['Title'])
```

```
[27]: merged_rating
```

```
[27]:
```

	Title	No. of Ratings	Avg. Ratings
0	\$1,000,000 Duck (1971)	37	3.027027
1	'Night Mother (1986)	70	3.371429
2	'Til There Was You (1997)	52	2.692308
3	'burbs, The (1989)	303	2.910891
4	...And Justice for All (1979)	199	3.713568
...	...	...	...
3701	Zed & Two Noughts, A (1985)	29	3.413793
3702	Zero Effect (1998)	301	3.750831
3703	Zero Kelvin (Kjrlighetens kjtere) (1995)	2	3.500000
3704	Zeus and Roxanne (1997)	23	2.521739
3705	eXistenZ (1999)	410	3.256098

[3706 rows x 3 columns]

```
[28]: # we put a threshold of minimum 1000 ratings for a movie for sorting the
      ↪average ratings so that movies with few but high ratings don't falsely
      ↪influence the result
mr=merged_rating[merged_rating['No. of Ratings'] > 1000]
mr
```

```
[28]:
```

	Title	No. of Ratings	\
16	2001: A Space Odyssey (1968)	1716	
43	Abyss, The (1989)	1715	
68	African Queen, The (1951)	1057	
80	Air Force One (1997)	1076	
84	Airplane! (1980)	1731	
...	...	...	

3635	Willy Wonka and the Chocolate Factory (1971)	1313
3655	Witness (1985)	1046
3656	Wizard of Oz, The (1939)	1718
3679	X-Men (2000)	1511
3693	Young Frankenstein (1974)	1193

	Avg. Ratings
16	4.068765
43	3.683965
68	4.251656
80	3.588290
84	3.971115
...	...
3635	3.861386
3655	3.996176
3656	4.247963
3679	3.820649
3693	4.250629

[207 rows x 3 columns]

```
[29]: top_rating = mr.sort_values('Avg. Ratings', ascending = False).head(25)
top_rating
```

```
[29]:
```


	Title	No. of Ratings \
2970	Shawshank Redemption, The (1994)	2227
1354	Godfather, The (1972)	2223
3504	Usual Suspects, The (1995)	1783
2901	Schindler's List (1993)	2304
2711	Raiders of the Lost Ark (1981)	2514
2738	Rear Window (1954)	1050
3153	Star Wars: Episode IV - A New Hope (1977)	2991
975	Dr. Strangelove or: How I Learned to Stop Worr...	1367
609	Casablanca (1942)	1669
3015	Sixth Sense, The (1999)	2459
2055	Maltese Falcon, The (1941)	1043
2452	One Flew Over the Cuckoo's Nest (1975)	1725
684	Citizen Kane (1941)	1116
2401	North by Northwest (1959)	1315
1355	Godfather: Part II, The (1974)	1692
2990	Silence of the Lambs, The (1991)	2578
665	Chinatown (1974)	1185
2894	Saving Private Ryan (1998)	2653
2217	Monty Python and the Holy Grail (1974)	1599
1923	Life Is Beautiful (La Vita bella) (1997)	1152
3178	Sting, The (1973)	1049
127	American Beauty (1999)	3428

2112	Matrix, The (1999)	2590
2654	Princess Bride, The (1987)	2318
461	Boat, The (Das Boot) (1981)	1001

	Avg. Ratings
2970	4.554558
1354	4.524966
3504	4.517106
2901	4.510417
2711	4.477725
2738	4.476190
3153	4.453694
975	4.449890
609	4.412822
3015	4.406263
2055	4.395973
2452	4.390725
684	4.388889
2401	4.384030
1355	4.357565
2990	4.351823
665	4.339241
2894	4.337354
2217	4.335210
1923	4.329861
3178	4.320305
127	4.317386
2112	4.315830
2654	4.303710
461	4.302697

```
[30]: ax = top_rating.plot.bar(x='Title', y='Avg. Ratings',rot=90, ylabel='Avg.
      ↳Ratings')
```





#### 0.4 The ratings for all the movies reviewed by a particular user of user id = 2696


```
[31]: user2696=Master_Data[Master_Data['UserID']==2696]
      user2696=user2696[['Title','Rating']]
```

```
[32]: user2696
```

```
[32]:
```

	Title	Rating
24345	Back to the Future (1985)	2
29848	E.T. the Extra-Terrestrial (1982)	3
244232	L.A. Confidential (1997)	4
250014	Lone Star (1996)	5
273633	JFK (1991)	1
277808	Talented Mr. Ripley, The (1999)	4
371178	Midnight in the Garden of Good and Evil (1997)	4
377250	Cop Land (1997)	3
598042	Palmetto (1998)	4
603189	Perfect Murder, A (1998)	4
609204	Game, The (1997)	4
611956	I Know What You Did Last Summer (1997)	2
612552	Devil's Advocate, The (1997)	4
613486	Psycho (1998)	4
616546	Wild Things (1998)	4
618708	Basic Instinct (1992)	4
621101	Lake Placid (1999)	1
689379	Shining, The (1980)	4
697451	I Still Know What You Did Last Summer (1998)	2
777089	Client, The (1994)	3

```
[33]: ax1 = user2696.plot.barh(x='Title', y='Rating',rot=0, ylabel='Ratings of User_
      ↪2696')
```



## 0.5 Feature Engineering:

```
[34]: # splitting the different genres in the column Genre
genre_split=data_movie['Genre'].str.split('|')
genre_split
```

```
[34]: 0      [Animation, Children's, Comedy]
      1      [Adventure, Children's, Fantasy]
      2      [Comedy, Romance]
      3      [Comedy, Drama]
      4      [Comedy]
      ...
      3878      [Comedy]
      3879      [Drama]
      3880      [Drama]
      3881      [Drama]
      3882      [Drama, Thriller]
      Name: Genre, Length: 3883, dtype: object
```

```
[35]: #separate column for each genre category with a one-hot encoding ( 1 and 0)
      ↳whether or not the movie belongs to that genre.
genre_sep=data_movie['Genre'].str.get_dummies()
genre_sep
```

```
[35]:      Action  Adventure  Animation  Children's  Comedy  Crime  Documentary  \
0           0           0           1           1           1           0           0
1           0           1           0           1           0           0           0
2           0           0           0           0           1           0           0
3           0           0           0           0           1           0           0
4           0           0           0           0           1           0           0
```

...	...	...	...	...	...	...	...	...	...
3878	0	0	0	0	0	1	0	0	0
3879	0	0	0	0	0	0	0	0	0
3880	0	0	0	0	0	0	0	0	0
3881	0	0	0	0	0	0	0	0	0
3882	0	0	0	0	0	0	0	0	0

	Drama	Fantasy	Film-Noir	Horror	Musical	Mystery	Romance	Sci-Fi	\
0	0	0	0	0	0	0	0	0	
1	0	1	0	0	0	0	0	0	
2	0	0	0	0	0	0	1	0	
3	1	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	

...	...	...	...	...	...	...	...	...	...
3878	0	0	0	0	0	0	0	0	0
3879	1	0	0	0	0	0	0	0	0
3880	1	0	0	0	0	0	0	0	0
3881	1	0	0	0	0	0	0	0	0
3882	1	0	0	0	0	0	0	0	0

	Thriller	War	Western
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

...	...	...	...
3878	0	0	0
3879	0	0	0
3880	0	0	0
3881	0	0	0
3882	1	0	0

[3883 rows x 18 columns]

```
[36]: # Finding out all the unique genres
print(genre_sep.columns)
print(len(genre_sep.columns))
```

```
Index(['Action', 'Adventure', 'Animation', 'Children's', 'Comedy', 'Crime',
      'Documentary', 'Drama', 'Fantasy', 'Film-Noir', 'Horror', 'Musical',
      'Mystery', 'Romance', 'Sci-Fi', 'Thriller', 'War', 'Western'],
      dtype='object')
18
```

```
[37]: # features affecting the ratings of any particular movie
# gender, age, occupation, genre
```

```
# we create a dataframe consisting of only those relevant features

Ratings_Data=pd.concat([Master_Data[['Gender', 'Age', 'Occupation', 'Rating']],genre_sep], axis=1)
# replace Male(M) by 1 and Female(F) by 0 in the Gender column
Ratings_Data['Gender'].replace({'M': 1, 'F': 0 }, inplace=True)
```

```
[38]: Ratings_Data.head(10)
```

```
[38]:
```

	Gender	Age	Occupation	Rating	Action	Adventure	Animation	Children's	\
0	0	1	10	5	0.0	0.0	1.0	1.0	
1	1	56	16	5	0.0	1.0	0.0	1.0	
2	1	25	12	4	0.0	0.0	0.0	0.0	
3	1	25	7	4	0.0	0.0	0.0	0.0	
4	1	50	1	5	0.0	0.0	0.0	0.0	
5	0	18	3	4	1.0	0.0	0.0	0.0	
6	1	1	10	5	0.0	0.0	0.0	0.0	
7	0	25	7	5	0.0	1.0	0.0	1.0	
8	0	25	1	3	1.0	0.0	0.0	0.0	
9	1	45	3	5	1.0	1.0	0.0	0.0	

	Comedy	Crime	...	Fantasy	Film-Noir	Horror	Musical	Mystery	Romance	\
0	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	...	1.0	0.0	0.0	0.0	0.0	0.0	
2	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	1.0	
3	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
5	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
6	1.0	0.0	...	0.0	0.0	0.0	0.0	0.0	1.0	
7	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
8	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
9	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	

	Sci-Fi	Thriller	War	Western
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
5	0.0	1.0	0.0	0.0
6	0.0	0.0	0.0	0.0
7	0.0	0.0	0.0	0.0
8	0.0	0.0	0.0	0.0
9	0.0	1.0	0.0	0.0

```
[10 rows x 22 columns]
```

```
[39]: # Find the correlation of the other features with Rating to see which is
      ↪maximum correlated
      Ratings_Data1=Ratings_Data[['Age','Occupation','Rating','Gender']]
      Ratings_Data1[Ratings_Data1.columns].corr()['Rating']
```

```
[39]: Age          0.056869
      Occupation   0.006753
      Rating       1.000000
      Gender      -0.019861
      Name: Rating, dtype: float64
```

## 0.6 An appropriate model to predict the movie ratings: Linear Regression

```
[40]: from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import train_test_split
      from sklearn import metrics
```

```
[41]: X_feature=Ratings_Data1.drop(['Rating'],axis=1)
      Y_target=Ratings_Data1['Rating']
```

```
[42]: print(X_feature.shape)
      X_feature.head
```

```
(1000209, 3)
```

```
[42]: <bound method NDFrame.head of          Age  Occupation  Gender
0          1          10        0
1         56          16        1
2         25          12        1
3         25           7        1
4         50           1        1
...      ...      ...      ...
1000204    18          17        1
1000205    35          14        1
1000206    18          17        1
1000207    18          20        0
1000208    25           1        1

[1000209 rows x 3 columns]>
```

```
[43]: print(Y_target.shape)
      Y_target.head
```

```
(1000209,)
```

```
[43]: <bound method NDFrame.head of 0          5
      1          5
      2          4
      3          4
      4          5
      ..
      1000204    5
      1000205    3
      1000206    1
      1000207    5
      1000208    4
      Name: Rating, Length: 1000209, dtype: int64>
```

```
[44]: x_train, x_test, y_train, y_test = train_test_split(X_feature, Y_target,
      ↪random_state=1)
```

```
[45]: print(x_train.shape)
      print(x_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(750156, 3)
(250053, 3)
(750156,)
(250053,)
```

```
[46]: lin_reg = LinearRegression()
```

```
[47]: lin_reg.fit(x_train, y_train)
```

```
[47]: LinearRegression()
```

```
[48]: y_pred = lin_reg.predict(x_test)
```

```
[49]: y_pred
```

```
[49]: array([3.55002686, 3.65282438, 3.51390249, ..., 3.59487431, 3.60326881,
      3.68434351])
```

```
[50]: # print the values obtained from different classification metrics
      print('y-intercept: ', lin_reg.intercept_)
      print('Beta coefficients: ',lin_reg.coef_)
      print('Mean Sq Error MSE: ',metrics.mean_squared_error(y_test, y_pred))
      print('Root Mean Sq Error RMSE:',np.sqrt(metrics.mean_squared_error(y_test,
      ↪y_pred)))
      print('r2 value: ',metrics.r2_score(y_test, y_pred))
```

```
y-intercept: 3.4518977639059805
Beta coefficients: [ 0.00540498  0.00083945 -0.05207392]
Mean Sq Error MSE: 1.2431690877572024
Root Mean Sq Error RMSE: 1.1149749269634732
r2 value: 0.0034613427979662825
```

```
[51]: # print the first 20 actual and predicted responses
print('actual: ', y_test.values[0:10])
print('predicted: ', y_pred[0:10])
```

```
actual:      [4 4 3 3 4 5 4 3 4 4]
predicted:   [3.55002686 3.65282438 3.51390249 3.70855308 3.58899816 3.50047129
 3.5357878 3.50047129 3.68434351 3.58899816]
```

```
[ ]:
```